



Marketing Science

Publication details, including instructions for authors and subscription information:
<http://pubsonline.informs.org>

Managing Churn to Maximize Profits

Aurélie Lemmens, Sunil Gupta

To cite this article:

Aurélie Lemmens, Sunil Gupta (2020) Managing Churn to Maximize Profits. Marketing Science 39(5):956-973. <https://doi.org/10.1287/mksc.2020.1229>

Full terms and conditions of use: <https://pubsonline.informs.org/Publications/Librarians-Portal/PubsOnLine-Terms-and-Conditions>

This article may be used only for the purposes of research, teaching, and/or private study. Commercial use or systematic downloading (by robots or other automatic processes) is prohibited without explicit Publisher approval, unless otherwise noted. For more information, contact permissions@informs.org.

The Publisher does not warrant or guarantee the article's accuracy, completeness, merchantability, fitness for a particular purpose, or non-infringement. Descriptions of, or references to, products or publications, or inclusion of an advertisement in this article, neither constitutes nor implies a guarantee, endorsement, or support of claims made of that product, publication, or service.

Copyright © 2020, INFORMS

Please scroll down for article—it is on subsequent pages



With 12,500 members from nearly 90 countries, INFORMS is the largest international association of operations research (O.R.) and analytics professionals and students. INFORMS provides unique networking and learning opportunities for individual professionals, and organizations of all types and sizes, to better understand and use O.R. and analytics tools and methods to transform strategic visions and achieve better outcomes.

For more information on INFORMS, its publications, membership, or meetings visit <http://www.informs.org>

Managing Churn to Maximize Profits

Aur lie Lemmens,^a Sunil Gupta^b

^aRotterdam School of Management, Erasmus University Rotterdam, Rotterdam, Netherlands; ^bHarvard Business School, Boston, Massachusetts 02163

Contact: lemmens@rsm.nl,  <https://orcid.org/0000-0002-7630-6159> (AL); sgupta@hbs.edu (SG)

Received: January 24, 2017

Revised: June 19, 2019; December 4, 2019

Accepted: December 13, 2019

Published Online in Articles in Advance:
August 27, 2020

<https://doi.org/10.1287/mksc.2020.1229>

Copyright:   2020 INFORMS

Abstract. Customer defection threatens many industries, prompting companies to deploy targeted, proactive customer retention programs and offers. A conventional approach has been to target customers either based on their predicted churn probability, or their responsiveness to a retention offer. However, both approaches ignore that some customers contribute more to the profitability of retention campaigns than others. This study addresses this problem by defining a profit-based loss function to predict, for each customer, the financial impact of a retention intervention. This profit-based loss function aligns the objective of the estimation algorithm with the managerial goal of maximizing the campaign profit. It ensures (1) that customers are ranked based on the incremental impact of the intervention on churn and postcampaign cash flows, after accounting for the cost of the intervention, and (2) that the model minimizes the cost of prediction errors by penalizing customers based on their expected profit lift. Finally, it provides a method to optimize the size of the retention campaign. Two field experiments affirm that this approach leads to significantly more profitable campaigns than competing models.

History: K. Sudhir served as the editor-in-chief and Scott Neslin served as associate editor for this article.

Funding: A. Lemmens was supported by Veni Personal Grant [451-09-005] and Vidi Personal Grant [452-12-011] from the Dutch National Science Foundation (NWO).

Supplemental Material: The data files and online appendices are available at <https://doi.org/10.1287/mksc.2020.1229>.

Keywords: defection • field experiments • loss function • machine learning • proactive churn management • profit lift • stochastic gradient boosting

1. Introduction

Customer defection is a global phenomenon, as exemplified by the estimated 20% annual churn rates for credit cards in the United States and 20%–38% annual churn rates for mobile phone carriers in Europe (Bobbier 2013). As customer acquisition costs continue to rise, managing customer churn has become critically important for the profitability of companies. According to McKinsey & Co., reducing churn could increase earnings of a typical U.S. wireless carrier by nearly 10% (Braff et al. 2003).

Not surprisingly, top executives cite customer retention as a top marketing priority, which they pursue with higher retention budgets (Forbes 2011) and more sophisticated, proactive churn management programs. These retention programs attempt to target potential churners with incentives (Ganesh et al. 2000), such as special offers, discounts, personalized (e)mail, or gifts, all of which aim to boost targeted customers' behavioral loyalty (Winer 2001).

For years, marketing practice and research have mainly focused on churn prediction and proposed ways to target customers according to their estimated churn risk (Lemmens and Croux 2006, Neslin et al. 2006, Risselada et al. 2010, Schweidel et al. 2011,

Ascarza and Hardie 2013; for a review, see Ascarza et al. 2018). Despite the popularity of this approach, recent studies have found that ranking customers on the basis of churn probability may lead to ineffective retention campaigns. Instead, Ascarza (2018) and Guelman et al. (2012) propose to rank customers based on their sensitivity to the intervention, regardless of their risk of churning. Using *uplift* random forests, they identify customers for whom the intervention will prompt the greatest lift in retention probability.

In both cases, however, existing approaches fail to recognize the ultimate goal of companies to maximize the *profit* of their proactive retention campaigns. First, the rankings provided by both approaches are solely based on churn (risk or lift) and ignore the profit impacts of a retention intervention. The *profit lift* of a proactive retention incentive can be estimated according to the intervention's (1) impact on the churn probability of the targeted customer, (2) impact on postcampaign cash flows, and (3) cost. A positive profit lift indicates that the intervention is likely to increase their retention probabilities and/or postcampaign revenues sufficiently to compensate for the intervention cost. In contrast, negative profit lifts signal cases in which targeting would lead to a loss for

the company. By focusing on churn lift rather than on profit lift, prior approaches might end up targeting customers for whom the effect on retention is positive, but the profit lift is negative.

Second, past approaches to obtain customer rankings ignore that prediction errors are more consequential in terms of campaign profit for some customers than others. Churn models attempt to minimize misclassifications of all customers' churn, regardless of their profit potential. Likewise, uplift models aim at estimating each customer's conditional average treatment effect as accurately as possible. In reality, however, not all customers are equally valuable, and incorrectly predicting the churn risk or lift would be costlier for some customers than others. For example, failing to predict the response of a high profit-lift customer will have a larger financial impact on the campaign profit than failing to predict the response of a customer who is insensitive to the intervention.

We propose a new approach that addresses these two limitations. Our approach defines a profit-based loss function that estimates, for every customer, the expected profit lift, taking into account the customer-specific cost of a prediction error. In contrast to existing approaches, the profit-based loss function fully aligns with the managerial objective of the retention campaign. The profit-based loss function weights more heavily those customers who have the greatest (positive or negative) impact on the campaign profit. This weighted estimation offers more accurate predictions for high profit-lift customers and thus boosts the profitability of the campaign. We empirically demonstrate the superiority of our approach for two retention campaigns: for a European interactive television subscription firm (Datta et al. 2015) and for a special interest membership organization in North America.

The remainder of this article is organized as follows. In Section 2, we review existing approaches to customer retention and explain how they differ from a profit-based estimation approach. In Section 3, as a key building block for our approach, we define the profit lift of a retention program. In Section 4, we construct the profit-based loss function. Section 5 describes the various steps of our approach. Section 6 outlines the alternative methods we use as benchmarks. In Section 7, we present our empirical applications and results. Finally, we conclude in Section 8 with some limitations and potential extensions.

2. Existing Approaches for Proactive Retention Management

Most customer retention research focuses on predicting churn. In a modeling tournament organized by the Teradata Center at Duke University, 44 academics and practitioners proposed models (for overviews,

see Neslin, et al. 2006, Blattberg et al. 2008), including logistic regression (Risselada et al. 2010, Knox and Van Oest 2014), probit models (Datta et al. 2015), Classification and Regression Trees or CART (Huang et al. 2012), neural nets (Huang et al. 2012), random forests (Larivière and Van den Poel 2005), bagging and boosting (Lemmens and Croux 2006), hazard models (Bolton 1998, Donkers, et al. 2007, Schweidel et al. 2008a, Braun and Schweidel 2011), hidden Markov models (Schweidel et al. 2011, Ascarza and Hardie 2013, Schweidel and Knox 2013), simultaneous equation models (Reinartz et al. 2005), probability models (Borle et al. 2008, Singh et al. 2009, Fader and Hardie 2010, Fader et al. 2010), and heuristics (Wübben and Von Wangenheim 2008).

Most of these articles focus on predicting churn, instead of estimating the *impact* of marketing interventions on churn. This perspective contrasts with other areas of marketing that focus on the incremental “effect of a marketing action to inform targeting decisions” (Ascarza 2018, p.82; see, e.g., Rossi et al. 1996, Venkatesan and Kumar 2004, Lewis 2005b, Venkatesan et al. 2007, Kumar et al. 2008, Khan et al. 2009, Neslin et al. 2009). Some notable exceptions include estimating the impacts of promotional activities (Schweidel et al. 2008b), retention dollars (Reinartz et al. 2005), or targeting the social network of an individual (Godinho de Matos et al. 2018) on churn and profit, but the effects are estimated at an *aggregate* or *segment*-specific level. Guelman et al. (2012), Datta et al. (2015), and Ascarza (2018) propose targeting customers according to their *individual* churn lift or change in churn caused by marketing intervention by estimating the heterogeneous treatment effect of the marketing action on churn.

Despite their differences (focus on churn risk or churn lift), all these approaches seek to minimize prediction errors for all customers regardless of their contribution to the campaign profit. Churn models aim to minimize the percentage of churners classified as nonchurners or vice versa; uplift models seek to reduce prediction errors in churn lift for all customers. The same limitation applies to studies that model customer churn and usage jointly to account for the dependence between these processes (Donkers et al. 2007, Borle et al. 2008, Ascarza and Hardie 2013, Datta et al. 2015), which is conceptually and mathematically different from penalizing the prediction errors of a model based on profit lift.

Empirical research thus tends to ignore the risk of using a loss function that is not aligned with managerial objectives. However, the loss function is integral to the model specification. It implicitly defines the model under consideration (Engle 1993) and should reflect the focal business problem (Granger 1969, Christoffersen and Jacobs 2004). When different

loss functions apply to in-sample estimations and out-of-sample evaluations, the mismatch leads to sub-optimal model selection and predictions (Engle 1993, Granger 1993). We note some exceptions: Using profit-based loss functions, Blattberg and George (1992) model customers' price sensitivity, Bult (1993) and Bult and Wittink (1996) estimate responses to mail, and Gladly et al. (2009) model temporal changes in usage. With conjoint analyses, Toubia and Hauser (2007) and Gilbride et al. (2008) identify managerially relevant loss functions. Bayesian statistics also highlight the importance of selecting managerially relevant priors (Montgomery and Rossi 1999). Data science and machine learning advances reiterate this importance (Chintagunta et al. 2016), such that firms such as Amazon.com seek to include managerially relevant loss functions in their data acquisition strategies (Saar-Tsechansky and Provost 2007). Surprisingly, this focus has been missing in the retention literature. Our proposal to define a profit-based loss function to estimate the profit lift of a retention intervention addresses this gap. Although conceptually straightforward, this approach requires significant changes in the model and estimation as indicated below.

3. Defining the Profit of Proactive Retention Actions

Imagine a proactive retention campaign that targets customers with a predefined retention incentive.¹ The firm's decision to target customer i is denoted T_i , so $T_i = 1$ indicates targeting, and $T_i = 0$ indicates no targeting. For every targeted customer i , the firm generates a *profit lift* π_i that represents the net impact of the intervention for this customer. In the potential outcome framework for causal inference (Rubin 2005), it corresponds to the conditional average treatment effect (CATE) of the retention program.² Following Neslin et al. (2006), the profit of the campaign is the sum of the profit lift of all targeted customers

$$\Pi = \sum_i^N \pi_i T_i, \quad (1)$$

where N is the total number of customers. In practice, we do not observe CATE because we do not know what the behavior of a targeted customer would have been if she was not targeted, nor the behavior of a nontargeted customer if she was targeted. Instead, we only observe one of both outcomes. One solution to estimate CATE is to run a randomized control trial, in which the intervention is randomized across a representative sample of customers. By observing the behavior of customers in both treatment and control groups, we can estimate the causal impact of the campaign at the customer level (Rosenbaum 2017).³

The expected profit lift of a retention action given the intervention cost δ is

$$E(\pi_i|\delta) = E(CLV_i - \delta|X_i, T_i = 1) - E(CLV_i|X_i, T_i = 0), \quad (2)$$

where $E(CLV_i - \delta|X_i, T_i = 1)$ is the net residual⁴ customer lifetime value (CLV) of customer i if targeted with an offer that costs δ , and $E(CLV_i|X_i, T_i = 0)$ is the (net) residual CLV if customer i is not targeted (Provost and Fawcett 2013). If a customer is targeted, her net residual CLV is the discounted value of the cash flows after the campaign minus the per customer cost of the retention intervention (Fader and Hardie 2010). When a customer is not targeted, the cost of the intervention is not incurred.

We consider the residual CLV for periods subsequent to the intervention (taking place at the beginning of period $t = 1$) in the general case where future retention probabilities and cash flows vary over time and given an infinite time horizon. Let $r_{it}^{(1)}$ and $r_{it}^{(0)}$ denote the retention probabilities of customer i in the period t following the intervention, depending on whether this customer is targeted (1) or not (0).⁵ Likewise, let $m_{it}^{(1)}$ and $m_{it}^{(0)}$ denote the cash flows generated by customer i in the period t following the intervention if targeted or not, conditional on customer i being alive. In addition, d is the discount rate for postcampaign cash flows. We distinguish between two types of retention incentives. Unconditional incentives, such as thank you presents, can be sent to customers without their prior consent and without any condition. Alternatively, conditional incentives (e.g., discounts, gifts) can be provided to customers only if they agree to purchase or renew their subscription. In many contractual settings, customers who are up for renewal receive a discount if they renew their contract. Depending on whether the retention incentive is conditional or not, we rewrite Equation (2) as

$$E(\pi_i|\delta) = \left(\sum_{t=1}^{\infty} \frac{r_{it}^{(1)} m_{it}^{(1)} - r_{it}^{(0)} m_{it}^{(0)}}{(1+d)^t} \right) - \frac{r_{i1}^{(1)} \delta}{(1+d)} \quad (3)$$

for unconditional incentives, and

$$E(\pi_i|\delta) = \left(\sum_{t=1}^{\infty} \frac{r_{it}^{(1)} m_{it}^{(1)} - r_{it}^{(0)} m_{it}^{(0)}}{(1+d)^t} \right) - \frac{\delta}{(1+d)} \quad (4)$$

for conditional ones.⁶ The difference between Equations (3) and (4) reflects that only customers who accept the offer prompt the cost of the conditional incentive. Also note that the overhead (fixed) costs of the retention campaign are not taken into account in

the profit lift because they do not affect the customer ranking. Finally, note that the churn lift definition provided by Ascarza (2018), given by $r_{i1}^{(1)} - r_{i1}^{(0)}$, is a special case of Equations (3) and (4) when the intervention cost and cash flows are ignored.

Theoretically, the net residual CLV should be estimated over an infinite time horizon, but, for practical purposes, most companies and academics focus on a specific time period and use a truncated CLV (Glady et al. 2015). In our empirical application, we estimate the impact of the intervention on the next period, as detailed in Section 7 (and further drop the time subscript in what follows). Indeed, estimating the *causal* effect of a retention intervention over an infinite, or even long, period of time is impractical because the company would need to ensure that no confounder influences the outcome of interest during this period. In practice, it is unlikely that the unconfoundedness assumption required for causal inference would not be violated (Rosenbaum and Rubin 1984).⁷

The expected profit lift can take any positive or negative value. It will be positive if the residual CLV, conditional on targeting, is larger than the combination of residual CLV in the absence of targeting and the cost of the retention intervention. For example, customers might intend to churn but change their mind after receiving the retention incentive or those who did not intend to churn might increase their spending in response to the incentive, because of *delight* (Blattberg et al. 2008). The expected profit lift instead is negative if the combination of the residual CLV in the absence of targeting and the cost of the intervention is greater than the residual CLV, conditional on targeting. Such counterproductive outcomes may occur if retention offers wake the sleeping dogs by reminding them of their dissatisfaction with the firm's service, thereby increasing their probability to churn (Ascarza et al. 2016).

4. Developing a Profit-Based Loss Function

In this section, we describe the classic loss function used in the domain of retention management. Based on the definition of profit lift proposed in the previous section, we then propose a new profit-based loss function. The profit-based loss function can be used with any estimation technique, including logistic regression (via likelihoods) and more advanced machine learning methods. For this study, we chose to rely on stochastic gradient boosting (SGB), a greedy algorithm based on gradient descent (Friedman 2001, 2002) because it supports flexible specifications of the loss function and provides powerful optimization.

4.1. Classic Loss Function

Let $(y_1, x_1), \dots, (y_i, x_i), \dots, (y_N, x_N)$ be a (calibration) sample of known values of y , the binary churn outcome, and x be a set of covariates for N customers. Let F be the function that maps x to y . The SGB estimation method we describe in the next section (or another binary prediction model, such as logistic regression) provides fitted values of $F(x_i)$ for every customer i based on the values of the x variables. When the fitted values are between 0 and 1, as with a logistic regression, they are called (churn) probabilities. When they are not restricted to this interval, as in the SGB method, they are called scores. Scores can be mapped to probabilities using various transformations⁸ (Greene 2003). For proactive retention programs, these estimated probabilities (or scores) represent the main input to rank-order customers, so companies can target the customers with the highest scores.

In a logistic regression, the estimation of probabilities relies on maximum likelihood, which aims to maximize the sum over the individual (weighted) log-likelihoods:

$$\log L_i = w_i \left(\tilde{y}_i \log p(x_i) + (1 - \tilde{y}_i) \log (1 - p(x_i)) \right), \quad (5)$$

where $\tilde{y}_i = 1$ when customer i is a churner or 0 when she is a nonchurner (Hastie et al. 2009). Most churn models assume a constant weight ($w_i = 1$) for all customers, so the cost of misclassification is the same for all of them. Some models add a customer-specific weight $w_i \neq 1$ (depending on the model used, weights must sum to one or not), which leads to a weighted estimator (Manski and Lerman 1977, Cosslett 1993). Weighted estimators can impose different costs on type I (false positives) and type II (false negatives) errors. This option is also available for imbalanced data (Lemmens and Croux 2006), such that different weights would be assigned to churners and nonchurners to account for the skewness of the y distribution.

Instead of maximizing a likelihood function, machine learning algorithms minimize a loss function. Most likelihoods have exact loss counterparts. Estimating a model with the log-likelihood in Equation (5) is the same as minimizing the binomial loss function,

$$\Psi_i = w_i \log \left(1 + e^{-2y_i F(x_i)} \right), \quad (6)$$

where $y_i = 1$ for a churner and -1 for a nonchurner, so $\tilde{y}_i = (y_i + 1)/2$ (for the proof, see Online Appendix A; Hastie et al. 2009). A loss function is defined by three components: its margin (here, $y_i F(x_i)$), its functional form, and, possibly, the weighing structure w_i . First, the margin defines the variable to predict (here, y) and qualifies the accuracy of a prediction of the outcome

of interest. In our preceding example, the margin $y_i F(x_i)$ captures the extent to which $F(x_i)$ is a good predictor of y_i . The more negative the margin becomes (i.e., y and F of opposite signs), the larger the prediction error is. The goal is to estimate a positive score F for $y_i = 1$ and a negative score F for $y_i = -1$.

Second, the functional form defines the loss assigned to a given observation, according to the estimation/prediction error associated with this observation. It indicates the predictions that need improvement. In our example, the loss is a monotone decreasing function of the margin $y_i F(x_i)$. The loss associated with a negative margin (i.e., higher error) is greater than that associated with a positive margin (smaller error). It only depends on y via the margin (i.e., in combination with F), so this loss cannot distinguish false positives ($y_i = -1$ and $F(x_i) > 0$) and false negatives ($y_i = 1$ and $F(x_i) < 0$) and instead penalizes both equally.

Third, the weight w_i determines an additional penalty assigned to a prediction error, specific to a given individual, similar to the weighted estimator we described previously. It can depend on y (penalize type I and type II errors differently) or other variables (e.g., cash flows generated by a customer). In most applications, each individual is weighted equally.

Although statistically sound, the likelihood and the loss function in Equations (5) and (6) do not align with the objectives of retention programs to maximize profits. They depend solely on whether a customer is well classified as a churner or not rather than on her profit lift.

4.2. Profit-Based Loss Function

In contrast with a classic loss function, a profit-based loss function seeks to ensure that the firm targets customers with a positive profit lift and does not target customers with a negative profit lift. To achieve these goals, we adapt the loss function in several ways.

First, we adapt the margin by replacing y_i with $E(\pi_i)$, to represent $E(\pi_i|\delta)$, as defined in Equation (2), which reflects the new outcome of interest. The new margin ensures that customers with a higher profit lift earn a higher score $F(x_i)$ than customers with a lower profit lift. The customer ranking based on these scores depends on the profit that a decision to target each of them would generate. The profit-based loss function thus becomes

$$\Psi_i = w_i \log(1 + e^{-2E(\pi_i)F(x_i)}). \quad (7)$$

Second, we weight the loss attached to each customer as a function of their expected profit lift to specify which prediction errors have the largest (positive or negative) impact on profit and thus should be penalized more. We empirically test three weighting schemes: (i) *symmetric weighting*, where $w_i = |E(\pi_i)|$

for all customers; (ii) *right weighting*, such that $w_i = |E(\pi_i)|$ for $E(\pi_i) \geq 1$ and $w_i = 1$ otherwise; and (iii) *left weighting*, where $w_i = |E(\pi_i)|$ for $E(\pi_i) \leq 1$ and $w_i = 1$ otherwise. Symmetric weighting ensures that predictions of the profit lift will be the most accurate for customers with the most extreme (positive or negative) profit lift values. Both ignoring a customer who would have contributed greatly to campaign success (i.e., $E(\pi_i)$ is much greater than 0) and mistakenly targeting a customer who reacts very negatively to the campaign (i.e., $E(\pi_i)$ is much smaller than 0) would have detrimental impacts in Equation (1), so symmetric weighting penalizes both equally. In contrast, right weighting focuses exclusively on customers with the most positive expected profit lifts, whereas left weighting focuses on customers with the greatest losses only. These latter weighting schemes mimic the notion of penalizing type I versus type II errors in classification settings.⁹ We treat the choice of the weighting scheme as an empirical question. The next section provides useful insight into the relative performance of these weighting schemes depending on the data characteristics.

4.3. Monte Carlo Simulations and Statistical Properties

The profit-based loss function belongs to the category of weighted estimators with endogenous weights (Solon et al. 2015). These estimators are known to have different statistical properties than unweighted estimators. In Online Appendix B, we report the results of two Monte Carlo simulations that study the statistical properties of the profit-based loss estimator. The simulations show why and when a weighted estimator outperforms an unweighted one. We find that an estimator that uses the profit-based loss function has the same statistical properties as the weighted estimator for endogenously stratified samples (Manski and Lerman 1977; Cosslett 1993; King and Zeng 2001a, b; Donkers et al. 2003). On average, these estimators are less efficient than estimators that use an unweighted loss, because weighing observations dilutes the information by assigning low weights to some observations.

However, the simulations reveal the mechanism by which a weighted estimator can offer more profitable campaigns than an unweighted one: The weighted estimator offers less bias and is more efficient than the unweighted estimator at the *individual* level for observations that receive a greater weight. In fact, only the predictions for the observations that receive a small weight exhibit a greater bias and lower efficiency. This cross-customer reallocation mechanism ensures that the weighted estimator fits the behavior of customers who have the greatest impact on the profit of the retention campaign better than the unweighted estimator does. We also find that this

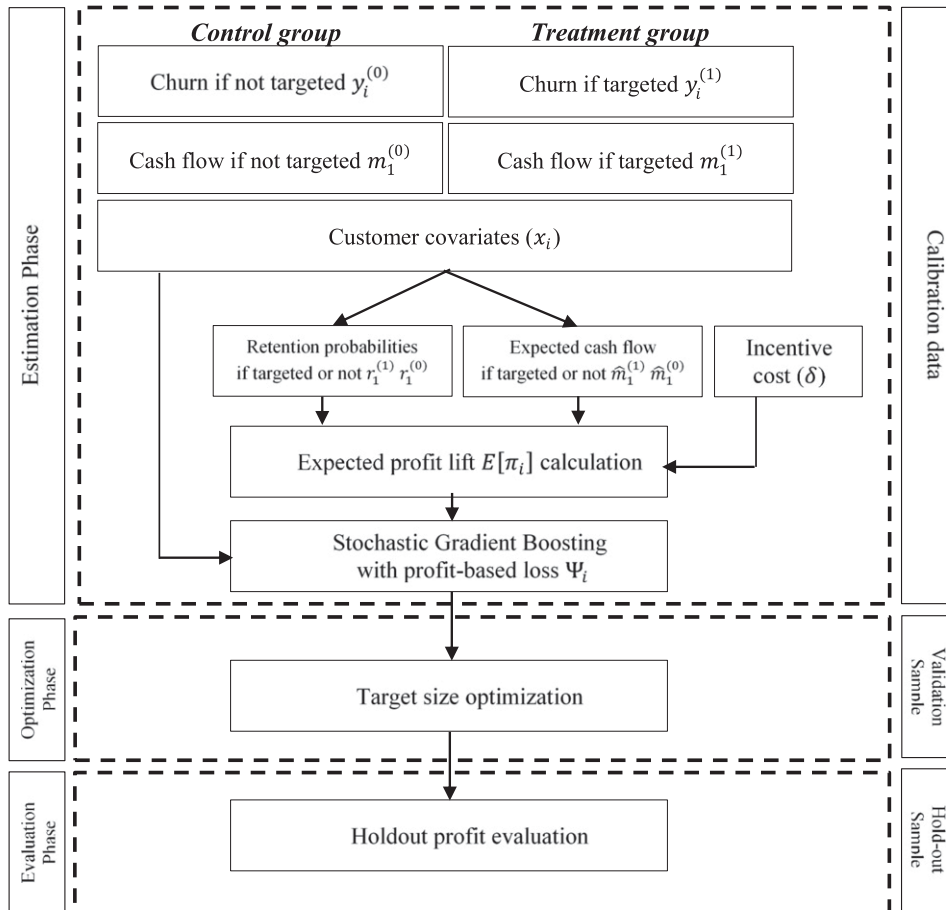
mechanism is stronger when the accuracy of statistical models is poorer (*smaller signal-to-noise ratio*), for *smaller sample sizes* and in presence of *less extreme weights*. Noting that retention models reputedly have low predictive power (see the recent review by Ascarza et al. 2018) and that field experiments are usually based on small treatment (and control) groups, we expect a substantial effect of weighting on the profitability of retention campaigns. Finally, the simulations show that the relative performance of the various weighting schemes (symmetric, right, and left weighting) likely depends on the expected profit lift distribution. In general, it is more beneficial to put greater weight on the under-represented part of the distribution (e.g., use right weighting if the share of positive profit lifts is small).

5. Integrated Profit-Based Approach

We integrate the profit-based loss function into the overall design of retention programs in three stages, as depicted in Figure 1: (1) estimate the expected profit lift of a retention intervention, (2) optimize the targets of the retention campaign, and (3) evaluate the

targeting decisions. Central to our approach, each stage uses a different sample of customers, which we refer to as the calibration sample, validation sample, and test sample, respectively. The calibration sample is used only for model estimation. We use the validation sample to determine the target size. The number of customers to target is chosen to maximize a holdout profit measure. As we further explain below, determining the target size based on a holdout profit measure allows us to account for the fact that our model might over- or underestimate the profit lift of customers on the calibration sample. Finally, the test sample contains customers who have not been used for estimation or for determining target size, so that we obtain a true holdout evaluation of campaign performance. To generate these three samples, we randomly split the data into three equal sets. To ensure the results are generalizable and not driven by any specific random split, we generated 100 different splits (Ascarza 2018) for both empirical applications. With this bootstrapping procedure, we also can test whether holdout performance is statistically superior to that of benchmark approaches.

Figure 1. Profit-Based Analysis Step by Step



5.1. Estimation Stage

We first estimate the heterogeneous treatment effect of the intervention on churn probabilities and cash flows separately, allowing for the possibility that the campaign can affect both processes differently. This step reflects the most recent benchmark in the literature (Guelman et al. 2012; Ascarza 2018, Hitsch and Misra 2018). Once we obtain an estimate of the components of the expected profit lift, we plug them into Equations (3) and (4), and we use SGB with the profit-based loss function defined in Equation (7). This step allows us to penalize customers according to their respective impact on campaign profitability.

5.1.1. Lift in Churn and Lift in Cash Flows. To estimate the heterogeneous treatment effect of an intervention, we use data from randomized experiments, in which some share of customers gets assigned to the treatment condition and targeted with a retention offer, whereas others are assigned to the control condition and do not receive a retention offer (in practice, this would mean running a pilot on a small sample of randomly chosen individuals). This approach is central to the potential outcome framework (Rubin 1974). Customers randomly allocated to two conditions should be, on average, similar in both observed and unobserved covariates across conditions, so we use these data to estimate the impact of the intervention. We observe churn in the period after intervention denoted by $y_i^{(1)}$ if the customer is in the treatment group and $y_i^{(0)}$ if the customer is in the control group. In addition, we observe customers' cash flows in the period following the company intervention: $m_i^{(1)}$ if the customer is in the treatment group and $m_i^{(0)}$ if in the control group. Finally, we observe customer-specific covariates and the cost of the retention offer (see Section 7). With this information, we estimate the customer-specific effect of the intervention on retention probabilities and cash flows.

Although estimating an average treatment effect of an intervention is straightforward (it only requires comparing the average outcome in the treatment and control groups), the estimation of the heterogeneous treatment effects is more complex because it requires comparing the outcomes for matched individuals. Machine learning for causal inference, and uplift models in particular, offers a solution to this problem by matching pairs of customers in the treatment and control groups on the basis of their available covariates and then comparing their respective churn and cash flow outcomes (Athey and Imbens 2016). Various uplift models can be used, depending on the nature of the dependent variable. For the binary retention model, we follow Ascarza (2018) and estimate the lift in retention probabilities $\hat{\tau}_i^{(0)}$ and $\hat{\tau}_i^{(1)}$ in the period following intervention using uplift random

forests (Guelman et al. 2015). To estimate the lift in cash flows $\hat{m}_i^{(0)}$ and $\hat{m}_i^{(1)}$ we use the uplift k -nearest neighbors (k NN) for continuous outcomes (Alemi et al. 2009, Su et al. 2012).¹⁰

5.1.2. Profit Lift Estimation by SGB with a Profit-Based Loss Function. We combine all estimates obtained in the previous step to calculate the expected profit lift according to Equations (3) or (4). Then, we estimate the profit lift scores with SGB using the profit-based loss function defined in Equation (7). This weighted loss function allows us to penalize customers according to their respective impact on campaign profitability. Any (machine learning) estimation method could be used with the profit-based loss function, but we choose SGB because of its superior predictive performance for churn prediction (it won the Teradata Churn modeling tournament; Lemmens and Croux 2006, Neslin et al. 2006) and other analyses (Hastie et al. 2009). Moreover, it uses a flexible optimization algorithm based on gradient descent, so it can be used with any loss function. As a greedy numerical optimization algorithm (Friedman et al. 2000, Friedman 2002), SGB sequentially combines predictions by simple models, typically regression trees (Breiman et al. 1983) and then makes initial guesses about each customer's outcome. It tries to predict residual errors by fitting a tree. At each iteration, a new tree is estimated to fit the residuals of the previous iteration. The estimation runs until no improvement occurs. We provide a description of regression trees in Online Appendix C.

Before the estimation, a loss function Ψ is chosen and used at each iteration b to compute the difference or error between the fitted scores $F_b(x_i)$, and the actual values to be predicted (in our case, the expected profit lifts). Once the loss function is defined, the estimation starts by setting each observation to an initial value, denoted by $\hat{F}_0(x_i)$, which can take any value in $(-\infty, \infty)$. From this initial guess, we compute the error (i.e., difference between the fitted values $\hat{F}_0(x_i)$ and actual values). The next step fits a tree model $T(x_i, \Theta_0)$ of the errors against the predictors x and computes the fitted values of these errors.¹¹ The number of terminal nodes is relatively small (maximum eight nodes) to avoid overfitting. These fitted errors are then combined with the predicted values $\hat{F}_0(x_i)$. The combination produces boosted fitted values (i.e., the original guess is *boosted* by the fitted errors), denoted $\hat{F}_1(x_i)$. This process repeats to compute the error from the boosted fitted values (difference between the fitted values $\hat{F}_1(x_i)$ and actual values), fit a tree model of the new errors, and combine the fitted values of these new errors to $\hat{F}_1(x_i)$. We repeat these steps B times until the model converges. Online Appendix D provides the estimation details.

5.2. Optimization Stage: Target Size Determination

The second stage determines how many and which customers to target to maximize the profit of the retention campaign. We use the validation sample for this purpose. The first step is to rank customers. Using the model estimated on the calibration sample, we predict the (holdout) profit lift scores for customers in the validation sample knowing their covariates values. We then rank them in order of decreasing scores, such that $\hat{F}(x_1) \geq \dots \geq \hat{F}(x_i) \geq \dots \geq \hat{F}(x_N)$. The final step is to determine the campaign size S , that is, the number of customers to target, starting at the top of the ranking. Because customers can have a negative profit lift, the optimal campaign size is usually smaller than 100%.

Two common approaches to determine target size include selecting the top decile (Lemmens and Croux 2006, Schweidel and Knox 2013) or applying a budget constraint (Datta et al. 2015), which we present in Section 6.2. Instead, we optimize the campaign size using full enumeration search, combined with offline evaluation. In particular, we calculate the holdout profit of a campaign of size S going from 1 to N (total number of customers in the validation sample) and identify the target size that maximizes the holdout campaign profit.

Calculating the holdout profit of a campaign of any target size S for a given customer ranking (i.e., the predictions of a model) is not straightforward, because we do not observe the actual profit lift of customers (i.e., we cannot observe the same unit at the same time in both the treatment and control groups). Offline policy evaluation provides a solution to this problem (Li et al. 2012). This evaluation strategy is common with randomized experiments (Ascarza 2018, Hitsch and Misra 2018). It is offline in the sense that it is not necessary to effectively target the customers identified by a given policy. Instead, analysts can leverage the random treatment allocation to test the performance of any policy. For each target size S , we estimate the impact of the campaign in the period following the intervention, according to the per customer profit lift π_S it generates. That is,

$$\pi_S = \frac{1}{N_t} \sum_{i \in \text{Treatment}} \left(m_i^{(1)} I(y_i^{(1)} = -1) - \delta \right) - \frac{1}{N_c} \sum_{j \in \text{Control}} m_j^{(0)} I(y_j^{(0)} = -1) \quad (8a)$$

or

$$\pi_S = \frac{1}{N_t} \sum_{i \in \text{Treatment}} m_i^{(1)} I(y_i^{(1)} = -1) - \frac{1}{N_c} \sum_{j \in \text{Control}} m_j^{(0)} I(y_j^{(0)} = -1) - \delta, \quad (8b)$$

depending on whether the offer is conditional (8a) or unconditional (8b), where N_t stands for the number of customers in the top S who actually received the retention incentive during the randomized experiment, and N_c is the number of customers in the top S who did not receive it. The first part indicates average postcampaign net cash flows (less the action cost) of customers in a target of size S that were effectively treated. The second part denotes the average postcampaign cash flows of customers in a target of size S that were not treated. The difference captures the *actual* per customer impact of the intervention on customers who belong to a target of size S . Note that π_S is an unbiased estimate of the actual profit of the targeting decisions conditional on S (Hitsch and Misra 2018). To obtain the total (holdout) profit of the designed retention campaign of size S , denoted Π_S , we multiply π_S by the number of customers targeted:

$$\Pi_S = S\pi_S. \quad (9)$$

Once we know Π_S for every target size S , we select the target size S^* that yields the highest holdout profit on the validation sample.¹²

5.3. Evaluation Phase: Holdout Profit of the Retention Campaign

Recall that S^* is determined with the validation sample, so strictly speaking, the value of Π_{S^*} on the validation sample is an in-sample measure of the campaign profit. We therefore use a third sample (test sample) to evaluate the holdout profit for a campaign of size S^* . The holdout profit of a campaign of target size S^* equals Π_{S^*} , calculated using Equation (9) for $S = S^*$ on the *test* sample.

6. Benchmark Models

We compare our approach against several benchmarks, including alternative estimation methods to rank-order customers and alternative approaches to determining target size.

6.1. Benchmark Estimation Methods to Rank Customers

6.1.1. Classic Loss. We estimate a churn model with SGB and the loss function defined in Equation (6) and rank customers on the basis of their estimated churn risk (Lemmens and Croux 2006).

6.1.2. Reordered Classic Loss. We reorder the classic loss scores (obtained from Section 6.1.1) by accounting for the profit that each customer is expected to generate if targeted with a retention action. Therefore, we predict the retention probabilities $\hat{r}_i^{(0)}$ and $\hat{r}_i^{(1)}$ using SGB with a classic loss function by setting the treatment dummy to 0 or 1, integrating the estimates

of postcampaign cash flows $\hat{m}_i^{(0)}$ and $\hat{m}_i^{(1)}$ (estimated with k -nearest neighbors; see Section 5.1.1), and finally plugging them into the profit lift formulas in Equations (3) or (4). This method is a so-called indirect (two-step) estimation approach (Hitsch and Misra 2018). This approach, or some version of it, is typically used by a vast majority of practitioners and scholars who are aware that ranking solely based on churn is flawed and that cash flows should also be taken into account.

6.1.3. Uplift Models. We estimate the lift in churn (Ascarza 2018) and the lift in cash flows as described in Section 5.1.1 and then combine them using Equations (3) or (4). This approach does not incorporate a profit-based loss function and corresponds to the first step of our estimation procedure.

6.2. Benchmark Methods to Determine the Optimal Target Size

6.2.1. Fixed Target Size. Companies often select target sizes by relying on managerial judgments or actual churn rates in their industry (i.e., higher churn rate prompts a larger target size). For this comparison, we define the target size according to the *churn rate* in our validation sample. Alternatively, we could determine the target size based on the available *budget* (Datta et al. 2015). Given the action cost, we calculate the number of customers that can be targeted with a specific budget. For illustration, we use a budget of 1,000 Euros or dollars.

6.2.2. Optimized Target Size Using Aggregate Metrics. Verbeke et al. (2012) calculate the optimal target size by combining information about the proportion of churners in the target, together with the average profit of targeting a customer. In their study, the probability of response to incentives and customer value used to calculate profit are hypothetical and assumed to be constant across all individuals. We extend this targeting rule by calculating the average treatment effect from the randomized controlled trial. It offers a crude approximation of the optimal target size selection we propose.

6.2.3. Buffer after Optimization. Finally, we consider the possibility that our target size optimization might be too restrictive and use an alternative target size that reflects a 10% buffer, such that it is larger than the optimized target size.¹³

7. Empirical Applications

We test our approach on two different customer databases from two different industries. The first data corresponds to an interactive television subscription service, provided by a firm located in continental

Europe and used by Datta et al. (2015).¹⁴ The second data refer to a subscription-based membership organization located in North America.

7.1. Interactive Television Subscription (Europe)

A major digital television provider in continental Europe offers access to local and international digital channels and video-on-demand (VOD) services. Customers pay for subscription that includes unlimited use of the basic iTV service (prices vary from approximately €20 to €100 per month depending on the type of service). In addition, customers can buy various additional packages (e.g., sports), for which they pay a higher monthly fee that varies across packages. Finally, customers can use the VOD rental service, for which they are charged on a pay-per-use basis, with an average price of €3 per VOD rental. To increase market penetration, the company offers new customers a free trial period of three months. About 40% of customers who use this service during the free trial period do not renew the service.

To decrease this high churn rate, the company used an intervention between August 2006 and July 2007. The retention offer was conditional on renewal, and its cost was about €12 per targeted customer, in line with practices in other industries.¹⁵ Not every subscriber ends the free trial period at the same time, so the intervention spanned nine waves, and during each wave, some proportion of free trial customers were targeted before the end of their trial period (treatment group), whereas others were not targeted (control group). We cannot identify the decision rules used to split the customers, so we used propensity score matching on the samples and performed a randomization check before and after the matching to ensure the final treatment and control groups are comparable¹⁶ (see Online Appendix E for details).

For each customer in the matched sample (2,595 treated and 2,595 not treated), we observe the month in which the retention offer was sent, whether the customer renewed the subscription after the intervention,¹⁷ the cash flows before and after the intervention, and other demographic and usage data (e.g., number of months the individual is a customer of the company, customer sex, age, language, household size, income based on zip code, installation method).

7.2. Special Interest Membership Organization (North America)

This special interest membership organization offers an annual membership for the right to use its services and receive discounts to attend events. The annual fee is approximately \$180. Each year, the organization sends out renewal letters to customers one month before their membership expires. The company ran a field experiment for five consecutive months that

tested whether adding a thank you gift to the renewal communication increased renewal rates. Each month, the company identified customers who were up for renewal and split them (randomly and evenly) between a treatment group that received a gift with the letter and a control group that received only the renewal letter. The per customer cost of the retention gift was about \$12. In total, 2,100 customers were involved in the experiment, and 1,044 of them were targeted. A randomization check confirms that randomization was done properly (Online Appendix E). This data set includes information on the month in which the renewal letter was sent, whether the customer renewed for the next year, and demographic and usage characteristics such as tenure (years), location (state where the member lives), whether the subscriber attended any organized or special interest event, and whether the subscriber had logged in to the organization's website.

7.3. Results

In the following sections, we compare the performance of a retention campaign with a profit-based loss function against the benchmarks from Section 6. We also explore the mechanisms that lead to the improved performance of our approach.

7.3.1. Financial Impact of Retention Campaign. In each bootstrap iteration, we apply the integrated profit-based approach described in Section 5 and calculate the corresponding holdout profit using Equations (8) and (9). Table 1 contains the average holdout profits over all bootstrapped samples for the classic loss function, reordered classic loss function, uplift model, and profit-based loss function (our approach). For our approach, we report the results for the weighting scheme that gives the highest performance. For study 1, all three weighting schemes give similar results (€4,967 for symmetric weighting, €4,945 for right weighting, and €5,026 for left weighting). For study 2, right weighting (\$1,328) significantly outperforms

symmetric weighting (\$413) and left weighting (minus \$55).¹⁸ We also report the bootstrapped differences between our approach and alternative ones, as well as the p values computed from the bootstrapped standard errors.

First, our approach leads to a more profitable retention campaign than all benchmark models. In both applications, the differences are highly significant. Note that this is the case for all three weighting schemes. In all three cases and both studies, our approach outperforms all other benchmarks. For example, in study 1, the campaign profit for our approach is 168% higher than for the classic loss, 300% higher than for the reordered classic loss, and 23% better than for the uplift method (which includes both uplift in churn and cash flow but does not use a profit-based loss function). In study 2, our approach is the only one that provides a positive campaign profit; all others lead to losses.


As expected, the classic loss function produces the lowest profit, because it is the only estimation method that optimizes a non-profit-related criterion. As we will show in a follow-up analysis, this method performs better when the assessment criterion reflects the optimization criterion (i.e., predicting churn rather than profit lift). By itself, this result confirms our main premise: Firms need to align their estimation method, and in particular their loss function, with their evaluation criterion, and both should fit their managerial objectives.

Second, reordering the scores to take the profit lift into account does not improve performance. Results for the classic loss and reordered classic loss are very similar to each other (as further illustrated by Figure 2, in which both curves follow very similar patterns). In general, the two-step reordering approach that minimizes the churn misclassification rate across all customers in the first step and then reorders customers according to their profit lift, performs significantly worse than our approach that incorporates customers' profit in the first step. This result corroborates findings by Hitsch and Misra (2018, p. 2),

Table 1. Average Campaign Holdout Profit for Different Estimation Methods (and Bootstrapped Differences with Profit-Based Loss)

Estimation method	Study 1: Interactive television subscription		Study 2: Special interest membership organization	
	Holdout profit (in Euro)	Difference (p value)	Holdout profit (in US\$)	Difference (p value)
Classic loss	€1,872.23	3,154 (0.000)	\$(1,669.37)	2,997 (0.000)
Reordered classic loss	€1,253.74	3,773 (0.000)	\$(1,709.16)	3,037 (0.000)
Uplift model	€4,092.97	933 (0.000)	\$(1,305.07)	2,633 (0.000)
Our approach	€5,026.36		\$1,327.76	

Notes. The last row, our approach, refers to the results provided by the profit-based loss function. The holdout profit column reports the average holdout profit obtained across all bootstrapped iterations. The difference column reports the bootstrapped difference between the holdout profit given by our approach and each alternative approach, together with the p values (in parentheses) obtained using the bootstrapped standard errors. All reported differences are significant at the 1% probability level.

Figure 2. Average Holdout Campaign Profit as a Function of Target Size for Different Estimation Methods

Notes. The curves represent the holdout profits of the campaign averaged over all bootstrap iterations. The horizontal dashed line represents a campaign targeting everyone.

who note that “methods that are trained to directly predict the incremental effect of targeting yield larger profits than conventional methods that indirectly predict the incremental effect based on the conditional expectation function that is trained on the outcome level.” For example, predicting churn (*outcome level*) for treated and control customers and then calculating the expected lift (*incremental effect*) is an indirect approach. Indirect approaches underperform direct approaches because the estimation uses the wrong metric. Reordering the ranking does not compensate for the loss function’s goal of minimizing errors in churn instead of profit lift.

Third, our approach significantly outperforms the uplift model that does not have profit-based loss function in estimation. The advantage of our approach over the uplift approach is that it *directly* estimates the profit lift (the profit lift is the dependent variable) and penalizes prediction errors for customers with the largest impact on campaign profits. As such, it aligns with firms’ managerial objectives of maximizing campaign profit.

Fourth, we note that in study 2, all approaches except our approach yield a negative total impact. The retention intervention was very ineffective in the first place and had a negative net impact on profits earned

from many customers. Despite this condition, the profit-based loss function can still identify a target size for which the total profit is positive.

7.3.2. Impact of Campaign Profits on Firm Revenue and Profit. Based on Table 1, we can assess the impact of our approach on the increase in firm profits from a proactive retention campaign. In study 1, our approach generates a per customer profit of €4.63 (€5,026/1,085 targeted customers). Given that the average annual revenue per customer in this data set is about €588, the profit earned from a proactive retention campaign would contribute about 1% to the firm’s annual revenue. This is substantial, considering that this profit results from a single campaign, captures its effect over a single period ahead, and its impact on firm profit (not revenue) would be even higher (e.g., in 2018, operating profit for Comcast in the United States was about 25% of revenue, which implies that a single retention campaign for the firm in our data set could increase its profits by about 4%). In study 2, our approach generates a per customer profit of \$4.99 (\$1,328/266 targeted customers). Given the annual subscription fee of \$180, this represents an increase of about 3% in annual revenue from a single campaign. Note, none of the other approaches

were able to achieve positive campaign profits. In summary, our approach has the potential to enhance a firm's future profits.

7.3.3. Profit as a Function of Campaign Size. Before we compare various target size optimization strategies, we explore how the profit of a campaign varies with its size by calculating a holdout cumulative profit for target sizes from 1% to 100% for each of the bootstrapped samples. Figure 2 reveals the average profit over all bootstrap samples for the four estimation methods.¹⁹ The profit with a 100% target size, when all customers receive a retention offer (i.e., it is the same across methods), indicates a positive impact in study 1 but a negative one in study 2. In the latter case, targeting all customers is not profitable.

This analysis confirms the superior performance of our approach and also reveals at least in part, why it works well. For both applications, customers who generate positive profits are ranked first. For study 1, the cumulative profit curve keeps increasing until it reaches its peak, at around 60% of the sample. Its slope is positive and larger than the slope of other methods, so our approach keeps adding profitable customers to the target, whereas other methods add less profitable or nonprofitable ones. In study 2, the overall poor impact of the intervention leads to a slightly different figure but similar conclusions. Our approach is the only one to achieve a positive campaign profit with small target sizes, because customers who contribute to the profit of the campaign are included first. The uplift model is the second best alternative, but it fails to rank high-profit customers first and thus requires a much larger target size—including a fair share of negative profit lift customers—before it reaches its maximum value (which is close to zero and far inferior to our approach). Finally, the profit curve for the classic loss function provides a good visualization of the problem of focusing on churn. Most customers who contribute to campaign profit are ranked low. For study 1, the slope of the classic loss is largest from 90% to 100%; for study 2, it is only positive from 70% to 80%. A different way of looking at it consists of decomposing, for each decile, the average treatment effect into the actual profit of a customer in the treatment group and the actual profit of a customer in the control group (see Online Appendix F).


7.3.4. The Role of Weighting. To complement these results and gain further insights into the role of the weighting scheme on the performance of our approach, we explore the relative performance of symmetric, right, and left weighting for both applications. Figure 3 reveals their average profit over all bootstrap samples as a function of campaign size.

Results are in line with those from the Monte Carlo simulation (Online Appendix B). The relative performance of the three weighting schemes depends on the distribution of the expected profit lift and confirms that it is more beneficial to put greater weight on the under-represented part of the distribution. For study 1, the share of customers with a negative expected profit lift (as inputted in the profit-based loss function) is slightly inferior to the proportion of customers with positive expected profit lifts (68% positive, 32% negative). As a result, left weighting slightly outperforms the other schemes (see Section 7.3.1). In contrast, customers with a positive expected profit lift (as inputted in the profit-based loss function) in study 2 are largely under-represented (17%), so right weighting is far more profitable than the other schemes. Figure 3 illustrates the function of the weighting scheme with regard to the performance of the prediction algorithm.

7.3.5. Determining Target Size. In Table 2, we compare the holdout profit of our target size optimization approach (Section 5.2) to the benchmarks (Section 6.2). For more details on the determination of the target size, we refer readers to Online Appendix G.

As expected, fixed target sizes determined before the estimation lead to significantly lower profits than our optimization strategy. Targeting as many customers as the number of expected churners or using a fixed budget is not a good strategy. Fixing the target size based on the churn rate offers 17% and 77% less profit than the optimized target size in study 1 and study 2. The fixed budget constraint leads to 92% and 77% less profit than the optimized target size in the two studies. In addition, our optimization strategy is superior to the optimization approach of Verbeke et al. (2012), which determines the target size on the basis of aggregate metrics that do not reflect customer heterogeneity in profit lift. Their approach yields 93% and 84% less profit than the optimized target size. Finally, in contrast with common practice, it is not preferable to add a buffer of customers to the retention campaign to ensure the target includes good ones. Doing so actually decreases the profit of the campaign by 3% and 47% for study 1 and study 2, respectively, because it adds nonprofitable customers.

7.3.6. Overlap of Customer Rankings. To understand the differences among various estimation methods, we investigate the extent to which the target identified by our approach overlaps with the targets of other estimation methods (see Ascarza 2018 for a similar approach). We rank-order customers according to the scores obtained by the various methods and split the four rankings into 10 deciles (the first decile corresponds to priority customers for targeting).

Figure 3. Average Holdout Campaign Profit as a Function of Target Size for Different Weighting Schemes

Notes. The curves represent the holdout profits of the campaign averaged over all bootstrap iterations. The horizontal dashed line represents a campaign targeting everyone.

For each decile, we then calculate the percentage overlap in customers targeted across methods. Figure 4 reveals the percentage overlap between our approach and all three alternatives. For instance, it indicates that almost 30% of the customers in the first decile provided by our (profit-based loss) approach also belong to the first decile provided by the uplift model (line with +).

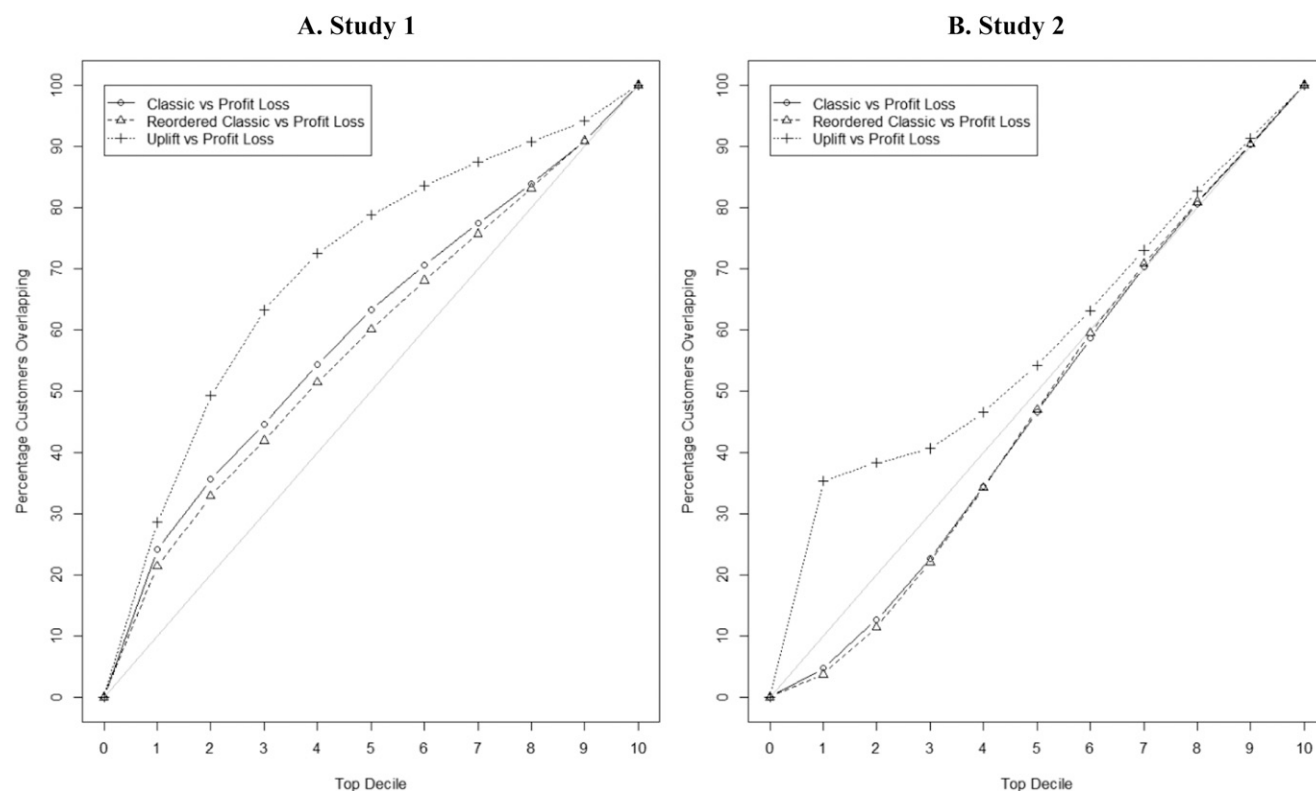
A value of 100% would mean that both groups perfectly overlap (i.e., the two approaches are identical in identifying profitable customers), whereas the 45-degree line represents a situation where the overlap between groups is purely because of chance. Figure 4 shows that the level of customer overlap between our approach (profit loss) and the classic loss or reordered

Table 2. Average Holdout Campaign Profit for Different Target Size Determination Methods (and Bootstrapped Differences with Profit-Based Loss)

Target size determination	Study 1: Interactive television subscription		Study 2: Special interest membership organization	
	Holdout profit (in Euro)	Difference (<i>p</i> value)	Holdout profit (in US\$)	Difference (<i>p</i> value)
Fixed target size				
Based on churn rate	€4,164.06	862 (0.000)	\$308.56	1,019 (0.000)
Based on budget	€398.76	4,628 (0.000)	\$304.98	1,023 (0.000)
Optimized target size				
Based on Verbeke	€345.46	4,681 (0.000)	\$207.94	1,120 (0.000)
10% buffer	€4,882.36	144 (0.001)	\$700.28	627 (0.010)
Our approach	€5,026.36		\$1,327.76	

Notes. The table provides the results of the profit-based loss function using the fixed and optimized target size selection approaches described in Section 6.2. The last row, our approach, refers to the proposed optimized target size selection using offline evaluation (Section 5.2). The holdout profit column reports the average holdout profit obtained across all bootstrapped iterations. The difference column reports the bootstrapped difference between the holdout profit given by our target size optimization approach and each alternative approach, together with *p* values (in parentheses) obtained using the bootstrapped standard errors. All reported differences are significant at the 1% probability level.

Figure 4. Average Percentage Customers Overlapping for Different Estimation Methods



Note. The 45% line corresponds to the level of overlap between two random rankings.

classic loss functions is close to random. In other words, our approach ranks customers very differently than these alternative rankings because they rely on different criteria. The greatest customer overlap is between our approach and the uplift model, but even this overlap is limited.

The higher profits obtained by our approach reflects contributions from customers who do not overlap, because this method identifies more profitable customers and excludes customers who have detrimental effects on total profits.

7.3.7. Drivers of Churn versus Drivers of Profit Lift. The differences between the profit-based and the classic loss functions translate into discrepancies in the factors explaining churn versus those that explain profit lift. We compute the relative variables' importance for both approaches (Friedman 2001). In study 1, household income contributes to almost 50% of the performance of the profit loss solution compared with only 11% for the classic loss solution after the treatment dummy (19%) and one of the sport package dummies (17%). In study 2, customers' geographical location (48% for the profit loss versus 25% for the classic loss) and tenure (28% for the profit loss versus 58% for the classic loss) are the

main drivers of profit lift and churn, but in different proportions.

7.3.8. Model Aligned with Managerial Objective. The additional profits earned with the profit-based loss function do not imply that it should be used in all circumstances. Table 3 reports the holdout Gini coefficient and top decile lifts for both studies, averaged across all bootstrap samples. The top decile lift measures the accuracy of the model in predicting churn among the top 10% riskiest customers. In turn, the Gini coefficient provides a measure of model accuracy in predicting churn for the entire customer base. The higher the top decile lift and Gini coefficient, the better the model predicts churn (see Lemmens and Croux 2006 for detailed definitions of these metrics).

The (reordered) classic loss function provides the most accurate predictions of customer churn behavior, because it is the only approach that seeks to minimize errors in churn prediction without considering other dimensions of customer profit. This shows that, if the goal is to accurately predict churn, the classic loss function is the best, but if the goal is to maximize retention campaign profit, our approach is more suitable.

Table 3. Average Holdout Churn Predictive Accuracy for Different Estimation Methods

Estimation method	Study 1: Interactive television subscription		Study 2: Special interest membership organization	
	Gini coefficient	Top decile lift	Gini coefficient	Top decile lift
Classic loss	0.277	2.014	0.102	1.183
Reordered classic loss	0.218	1.820	0.110	1.215
Uplift model	0.150	1.435	−0.027	0.011
Our approach	0.142	1.291	−0.016	0.084

Notes. The table reports the average Gini coefficients and top decile lifts obtained across all bootstrapped iterations. The last row, our approach, refers to the results provided by the profit-based loss function.

7.3.9. Replication for Other Estimation Methods.

Finally, we replicate our results for a different estimation approach. Namely, we compare the performance of a logistic regression (classic loss) and a weighted regression (profit loss) with the weights defined in Equations (3) and (4). The profit loss offers significantly ($p < 0.01$) larger holdout profits than the classic loss in both studies. In study 1, €3,331 for the classic loss and €4,907 for the profit loss; in study 2, minus \$1,112 for the classic loss and \$416 for the profit loss. This replication confirms that the loss function's choice drives the improvement in profits regardless of the estimation approach.

8. Conclusion, Limitations, and Further Research

We propose a method to optimize the profit of proactive retention campaigns. Our approach defines the profit lift of a retention intervention according to the potential outcome framework for causal inference (Rubin 2005). We demonstrate the benefits of using a profit-based loss function in estimating the financial impact of a targeted marketing intervention. Our findings highlight the need for marketing academics and practitioners to pay attention to the choice of loss function, a feature that is often neglected in model estimation processes. In particular, this choice should match managers' objectives.

Our approach potentially fits many contexts, within and outside marketing, where organizations seek to target a set of individuals with a specific intervention (e.g., catalog, mail, charitable giving, and personalized promotions). Estimating heterogeneous treatment effects is an exciting topic, featured in studies across economics and econometrics (Imbens and Rubin 2015), management (Godinho de Matos et al. 2017), and computer science (Pearl and Mackenzie 2018). For each application, it is critical to carefully determine the appropriate loss function. When building their own goal-oriented loss functions, decision makers should (1) ensure that the margin specifies the true outcome of interest (i.e., goal of the intervention)

and (2) use a weighting scheme that prioritizes customers who have the largest impact on the success of the intervention. This is relevant even in nonprofit contexts, such as for predicting patient compliance with medical treatments. In this case, the loss function could incorporate patient-specific health risks and benefits associated with complying with the medical treatment.

Our results also show that the optimization of the target size of a retention campaign has a significant impact on profits. Retention literature is surprisingly silent on this topic; it mostly focuses on ranking customers. We find that selecting a target size that maximizes the campaign profit leads to significantly more profitable campaigns than using the common rules. We thus hope managers attend to not only the estimation method used to rank customers but also the number of customers to target.

Several limitations of this paper offer fruitful research opportunities. First, our approach can rank-order customers according to the profit lift they produce, in response to a specific retention campaign. Both field experiments reflect a single, specific retention incentive, in line with recent attempts to estimate heterogeneous treatment effect models for customer relationship management (Provost and Fawcett 2013, Ascarza 2018, Hitsch and Misra 2018). An interesting further challenge would be to explore variations in customer responses depending on the type and depth of retention interventions and then determine the costs at which each response is maximized (Venkatesan et al. 2007). Firms might estimate consumers' profit lift distributions for various costs by testing various retention incentives and then use these estimates to determine the optimal intervention per customer.

Second, our approach does not consider the long-run impact of retention interventions. Assumptions of unconfoundness make it difficult and impractical to estimate the profit lift of a single intervention over a long period of time. Managers are often unwilling to isolate a group of customers from any marketing

intervention for a long time period. Concerns about legal customer privacy protections, which mandate that companies may only keep customer data for the shortest amount of time possible, also complicate experiments that run for long periods.²⁰

Third, we do not model whether a customer's sensitivity to an intervention depends on expectations of retention offers in the future (Lewis 2005a). In digital, connected economies, customers are more aware of the attractive discounts that others receive when they indicate an intention to churn. This phenomenon of strategic churning is an interesting area for research and could be captured using dynamic structural models (Khan et al. 2009).

We hope our work will foster more research in the area of predicting the individual treatment effects and remind the reader of the importance of aligning the loss function used for model estimation with the managerial objectives of the campaign.

Acknowledgments

The authors thank Eva Ascarza (Harvard), Christophe Croux (EDHEC Business School), Hannes Datta (Tilburg University), Bram Foubert (Maastricht), and Jason Roos (Rotterdam School of Management) for tremendous support.

Endnotes

¹ We assume a constant, exogenously determined retention offer and optimize the customer target for a specific intervention. In the last section, we briefly discuss how to generalize our approach for multiple offers.

² Some companies call this construct Delta CLV, suggesting a comparison of the customer lifetime value (CLV) of a customer if targeted versus not targeted. We prefer the term profit lift, to acknowledge that we incorporate the cost of the intervention.

³ We take the viewpoint of the firm and define the treatment as the firm sending a retention incentive to a customer (as proposed by Ascarza 2018). Thus, the treatment (TE) and the intent-to-treat (ITT) effects coincide.

⁴ By using the residual CLV, we ignore transactions and costs (including acquisition cost) that precede the campaign as they are irrelevant to the current campaign.

⁵ The retention probability $r_{it} = \prod_{k=1}^t \tilde{r}_{ik}$ with \tilde{r}_{ik} the retention probability going from period $k - 1$ to period k . For instance, the retention probability two periods after intervention equals to the product of the retention probability in period 1 (right after intervention) and the retention probability in period 2.

⁶ We assume that the cost of the incentive is incurred in the same period as the first postintervention cash flow is received, and thus discount its cost by one period.

⁷ Our discussions with several customer retention managers affirmed that the main barrier to using A/B testing is that they do not want to isolate groups of customers for a long time. Practical constraints make it impossible to exclude the risk of contamination by post-treatment marketing interventions or external factors that have nothing to do with the retention campaign but that endanger the comparability of the treatment and control groups. A multiperiod horizon thus would

be practically intractable, as is also the case for the data sets in our empirical application.

⁸ The estimated scores \hat{F}_i between $]-\infty, +\infty[$ can be transformed into deflection probability estimates \hat{p}_i between $[0,1]$ (e.g., when computing CLV) using the logistic (inverted-logit) formula, $\hat{p}_i = \frac{1}{1+\exp(-2\hat{F}_i)}$.

⁹ We thank the review team for this suggestion.

¹⁰ Uplift kNN computes the Euclidean distance between every pair of observations using all observed characteristics. Next, it selects k (here, $k = 1$) observation(s) in the calibration sample that is/are the closest to each observation in the validation (and test) sample and that received the treatment; $\hat{m}_i^{(1)}$ is the (average) cash flow for this/these k nearest neighbor(s). The same thing is done for the control sample. Both uplift random forests and k nearest neighbors are implemented in the uplift R package (Guelman 2014).

¹¹ Fitting the errors gradually forces the model to predict the residual variance in the dependent variable that was unexplained in the prior iteration. Thus, the estimation progressively concentrates on customers whose behavior is difficult to predict.

¹² An alternative to using offline policy evaluation would be to let S^* equal the number of customers in the validation sample whose profit lifts are predicted positive by the estimation algorithm (here, SGB). However, such an approach would be sensitive to the scale of the predicted scores. It would overestimate (underestimate) the target size if the model overestimates (underestimates) the number of customers with a positive profit lift. Instead, offline evaluation is insensitive to the scale of the predicted scores. The scores only serve to determine the ranking of customers, whereas the optimal target size is determined based on the *holdout* campaign profit. This approach is particularly useful when the predicted scores are not scaled (units have no meaning) as is the case with most machine learning methods.

¹³ We thank a reviewer for this suggestion.

¹⁴ We are extremely grateful to the authors of this paper for sharing their data with us.

¹⁵ Discussions with managers responsible for proactive retention programs confirm similar numbers. For example, a North European telecom firm cited an average cost of approximately 7 euros per customer.

¹⁶ The management team has changed since 2007, so we cannot specify the decision rules used previously. However, considering the vast customer data and the flexibility of our matching algorithm, we believe matching can capture them relatively accurately. The post-matching randomization check confirms that the matched treatment and control groups do not differ.

¹⁷ We use two operationalizations to measure churn and cash flow after intervention: (1) one month after the intervention and (2) three months after the intervention. In the empirical section, we present the results for one month, but using three months does not affect our conclusions.

¹⁸ See Section 7.3.4 for more details on the relative performance of the various weighting schemes.

¹⁹ Note that the profit figures in Table 1 do not directly correspond to Figure 2. The former figures are obtained by determining the optimal target size per bootstrap iteration such that we obtain 100 holdout campaign profit measures. This approach allowed us to test whether two approaches perform significantly differently from each other. In contrast, Figure 2 does not fix the target size but averages the profit curves obtained at each iteration. The latter is useful to see how the profit evolves with the campaign size.

²⁰ See https://ec.europa.eu/info/law/law-topic/data-protection/reform/rules-business-and-organisations/principles-gdpr/how-long-can-data-be-kept-and-it-necessary-update-it_en.

References

- Alemi F, Erdman H, Griva I, Evans CH (2009) Improved statistical methods are needed to advance personalized medicine. *Open Translational Med. J.* 1:16–20.
- Ascarza E (2018) Retention futility: Targeting high-risk customers might be ineffective. *J. Marketing Res.* 55(1):80–98.
- Ascarza E, Hardie BGS (2013) A joint model of usage and churn in contractual settings. *Marketing Sci.* 32(4):570–590.
- Ascarza E, Iyengar R, Schleicher M (2016) The perils of proactive churn prevention using plan recommendations: Evidence from a field experiment. *J. Marketing Res.* 53(1):46–60.
- Ascarza E, Neslin SA, Netzer O, Anderson Z, Fader PS, Gupta S, Hardie BGS, Lemmens A, Libai B, Neal D, Provost F, Schrift R (2018) In pursuit of enhanced customer retention management: Review, key issues, and future directions. *Customers Needs Solutions* 5(1–2):65–81.
- Athey S, Imbens G (2016) Recursive partitioning for heterogeneous causal effects. *Proc. National Acad. Sci. USA* 113(27):7353–7360.
- Blattberg RC, George EI (1992) Estimation under profit-driven loss functions. *J. Bus. Econom. Statist.* 10(4):437–444.
- Blattberg RC, Kim B-D, Neslin SA (2008) *Database Marketing: Analyzing and Managing Customers* (Springer, New York).
- Bobbier T (2013) Keeping the customer satisfied: The dynamics of customer defection, and the changing role of the loss adjuster. Report, The Chartered Institute of Loss Adjusters, London.
- Bolton RN (1998) A dynamic model of the duration of the customer's relationship with a continuous service provider: The role of satisfaction. *Marketing Sci.* 17(1):45–65.
- Borle S, Singh SS, Jain DC (2008) Customer lifetime value measurement. *Management Sci.* 54(1):100–112.
- Braff A, Passmore WJ, Simpson M (2003) Going the distance with telecom customers. *McKinsey Quart.* 4:82–93.
- Braun M, Schweidel DA (2011) Modeling customer lifetimes with multiple causes of churn. *Marketing Sci.* 30(5):881–902.
- Breiman L, Friedman J, Olshen R, Stone C (1983) *Classification and Regression Trees* (Wadsworth Publishing Chapman and Hall/CRC, Boca Raton, FL).
- Bult JR (1993) Semiparametric vs. parametric classification models: An application to direct marketing. *J. Marketing Res.* 30(3):380–390.
- Bult JR, Wittink DR (1996) Estimating and validating asymmetric heterogeneous loss functions applied to healthcare fund raising. *Internat. J. Res. Marketing* 13(3):215–226.
- Chintagunta P, Hanssens DM, Hauser JR (2016) Editorial—Marketing science and big data. *Marketing Sci.* 35(3):341–342.
- Christoffersen P, Jacobs K (2004) The importance of the loss function in option valuation. *J. Financial Econom.* 72(2):291–318.
- Cosslett SR (1993) Estimation from endogenously stratified samples. Maddala GS, Rao CR, Vinod HD eds. *Handbook of Statistics* (Elsevier, Amsterdam), 1–43.
- Datta H, Foubert B, Van Heerde HJ (2015) The challenge of retaining customers acquired with free trials. *J. Marketing Res.* 52(2):217–234.
- Donkers B, Franses PH, Verhoef PC (2003) Selective sampling for binary choice models. *J. Marketing Res.* 40(4):492–497.
- Donkers B, Verhoef PC, de Jong MG (2007) Modeling CLV: A test of competing models in the insurance industry. *Quantitative Marketing Econom.* 5(2):163–190.
- Engle RF (1993) On the limitations of comparing mean square forecast errors [comment]. *J. Forecasting* 12(8):642–644.
- Fader PS, Hardie BGS (2010) Customer-base valuation in a contractual setting: The perils of ignoring heterogeneity. *Marketing Sci.* 29(1):85–93.
- Fader PS, Hardie BGS, Shang J (2010) Customer-base analysis in a discrete-time noncontractual setting. *Marketing Sci.* 29(6):1086–1108.
- Forbes (2011) Bringing 20/20 foresight to marketing: CMOs seek a clearer picture of the customer. Accessed April 2020, https://www.forbes.com/forbesinsights/bringing_foresight_to_marketing/index.html.
- Friedman J, Hastie T, Tibshirani R (2000) Additive logistic regression: A statistical view of boosting. *Ann. Statist.* 28(2):337–407.
- Friedman JH (2001) Greedy function approximation: A gradient boosting machine. *Ann. Statist.* 29(5):1189–1232.
- Friedman JH (2002) Stochastic gradient boosting. *Comput. Statist. Data Anal.* 38(4):367–378.
- Ganesh J, Arnold MJ, Reynolds KE (2000) Understanding the customer base of service providers: An examination of the differences between switchers and stayers. *J. Marketing* 64(3):65–87.
- Gilbride TJ, Lenk PJ, Brazell JD (2008) Market share constraints and the loss function in choice-based conjoint analysis. *Marketing Sci.* 27(6):995–1011.
- Gladly N, Baesens B, Croux C (2009) Modeling churn using customer lifetime value. *Eur. J. Oper. Res.* 197(1):402–411.
- Gladly N, Lemmens A, Croux C (2015) Unveiling the relationship between the transaction timing, spending and dropout behavior of customers. *Internat. J. Res. Marketing* 32(1):78–93.
- Godinho de Matos M, Ferreira P, Belo R (2018) Target the ego or target the group: Evidence from a randomized experiment in proactive churn management. *Marketing Sci.* 37(5):793–811.
- Godinho de Matos M, Ferreira P, Smith MD (2017) The effect of subscription video-on-demand on piracy: Evidence from a household-level randomized experiment. *Management Sci.* 64(12):5610–5630.
- Granger CWJ (1969) Prediction with a generalized cost of error function. *J. Oper. Res. Soc.* 20(2):199–207.
- Granger CWJ (1993) On the limitations of comparing mean square forecast errors [comment]. *J. Forecasting* 12(8):651–652.
- Greene WH (2003) *Econometric Analysis*, 6th ed. (Prentice Hall, Upper Saddle River, NJ).
- Guelman L (2014) Uplift: Uplift modeling. R package version 0.3.5. Accessed April 2020, <https://CRAN.R-project.org/package=uplift>.
- Guelman L, Guillén M, Pérez-Marín AM (2012) Random forests for uplift modeling: An insurance customer retention case. Engemann KJ, Gil-Lafuente AM, Merigó JM, eds. *Modeling and Simulation in Engineering, Economics, and Management* (Springer, New York), 123–133.
- Guelman L, Guillén M, Pérez-Marín AM (2015) Uplift random forests. *Cybernet Systems* 46(3):230–248.
- Hastie T, Tibshirani R, Friedman JH (2009) *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. (Springer, New York).
- Hitsch GJ, Misra S (2018) Heterogeneous treatment effects and optimal targeting policy evaluation. Preprint, submitted February 6, <http://dx.doi.org/10.2139/ssrn.3111957>.
- Huang B, Kechadi MT, Buckley B (2012) Customer churn prediction in telecommunications. *Expert Systems Appl.* 39(1):1414–1425.
- Imbens GW, Rubin DB (2015) *Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction* (Cambridge University Press, Cambridge, UK).
- Khan R, Lewis M, Singh V (2009) Dynamic customer management and the value of one-to-one marketing. *Marketing Sci.* 28(6):1063–1079.
- King G, Zeng L (2001a) Explaining rare events in international relations. *Internat. Organ.* 55(3):693–715.
- King G, Zeng L (2001b) Logistic regression in rare events data. *Political Anal.* 9(2):137–163.
- Knox G, Van Oest R (2014) Customer complaints and recovery effectiveness: A customer base approach. *J. Marketing* 78(5):42–57.
- Kumar V, Venkatesan R, Bohling T, Beckmann D (2008) Practice Prize Report—The power of CLV: Managing customer lifetime value at IBM. *Marketing Sci.* 27(4):585–599.

- Larivière B, Van den Poel D (2005) Predicting customer retention and profitability by using random forests and regression forests techniques. *Expert Systems Appl.* 29(2):472–484.
- Lemmens A, Croux C (2006) Bagging and boosting classification trees to predict churn. *J. Marketing Res.* 43(2):276–286.
- Lewis M (2005a) Incorporating strategic consumer behavior into customer valuation. *J. Marketing* 69(4):230–238.
- Lewis M (2005b) Research note: A dynamic programming approach to customer relationship pricing. *Management Sci.* 51(6):986–994.
- Li L, Chu W, Langford J, Moon T, Wang X (2012) An unbiased offline evaluation of contextual bandit algorithms with generalized linear models. Glowacka D, Dorard L, Shawe-Taylor J, eds. *Proc. Workshop Online Trading Exploration Exploitation 2 (JMLR: Workshop and Conference Proceedings, Bellevue, Washington)* 26:19–36.
- Manski CF, Lerman SR (1977) The estimation of choice probabilities from choice based samples. *Econometrica* 45(8):1977–1988.
- Montgomery AL, Rossi PE (1999) Estimating price elasticities with theory-based priors. *J. Marketing Res.* 36(4):413–423.
- Neslin SA, Gupta S, Kamakura W, Lu J, Mason CH (2006) Defection detection: Measuring and understanding the predictive accuracy of customer churn models. *J. Marketing Res.* 43(2):204–211.
- Neslin SA, Novak TP, Baker KR, Hoffman DL (2009) An optimal contact model for maximizing online panel response rates. *Management Sci.* 55(5):727–737.
- Pearl J, Mackenzie D (2018) *The Book of Why: The New Science of Cause and Effect* (Basic Books, New York).
- Provost F, Fawcett T (2013) *Data Science for Business: What You Need to Know About Data Mining and Data-Analytic Thinking* (O'Reilly Media, Inc., Sebastopol, CA).
- Reinartz W, Thomas JS, Kumar V (2005) Balancing acquisition and retention resources to maximize customer profitability. *J. Marketing* 69(1):63–79.
- Risselada H, Verhoef PC, Bijmolt THA (2010) Staying power of churn prediction models. *J. Interactive Marketing* 24(3):198–208.
- Rosenbaum PR, Rubin DB (1984) Reducing bias in observational studies using subclassification on the propensity score. *J. Amer. Statist. Assoc.* 79(387):516–524.
- Rosenbaum PR (2017) *Observation and Experiment: An Introduction to Causal Inference* (Harvard University Press, Cambridge, MA).
- Rossi PE, McCulloch RE, Allenby GM (1996) The value of purchase history data in target marketing. *Marketing Sci.* 15(4):321–340.
- Rubin DB (1974) Estimating causal effects of treatments in randomized and nonrandomized studies. *J. Educ. Psych.* 66(5):688–701.
- Rubin DB (2005) Causal inference using potential outcomes: Design, modeling, decisions. *J. Amer. Statist. Assoc.* 100(469):322–331.
- Saar-Tsechansky M, Provost F (2007) Decision-centric active learning of binary-outcome models. *Inform. System Res.* 18(1):4–22.
- Schweidel DA, Knox G (2013) Incorporating direct marketing activity into latent attrition models. *Marketing Sci.* 32(3):471–487.
- Schweidel DA, Bradlow ET, Fader PS (2011) Portfolio dynamics for customers of a multiservice provider. *Management Sci.* 57(3):471–486.
- Schweidel DA, Fader PS, Bradlow ET (2008a) A bivariate timing model of customer acquisition and retention. *Marketing Sci.* 27(5):829–843.
- Schweidel DA, Fader PS, Bradlow ET (2008b) Understanding service retention within and across cohorts using limited information. *J. Marketing* 72(1):82–94.
- Singh SS, Borle S, Jain DC (2009) A generalized framework for estimating customer lifetime value when customer lifetimes are not observed. *Quantitative Marketing Econom.* 7(2):181–205.
- Solon G, Haider SJ, Wooldridge JM (2015) What are we weighting for? *J. Human Resources* 50(2):301–316.
- Su X, Kang J, Fan J, Levine RA, Yan X (2012) Facilitating score and causal inference trees for large observational studies. *J. Machine Learn. Res.* 13(10):2955–2994.
- Toubia O, Hauser JR (2007) Research note—On managerially efficient experimental designs. *Marketing Sci.* 26(6):851–858.
- Venkatesan R, Kumar V (2004) A customer lifetime value framework for customer selection and resource allocation strategy. *J. Marketing* 68(4):106–125.
- Venkatesan R, Kumar V, Bohling T (2007) Optimal customer relationship management using Bayesian decision theory: An application for customer selection. *J. Marketing Res.* 44(4):579–594.
- Verbeke W, Dejaeger K, Martens D, Hur J, Baesens B (2012) New insights into churn prediction in the telecommunication sector: A profit driven data mining approach. *Eur. J. Oper. Res.* 218(1):211–229.
- Winer RS (2001) A framework for customer relationship management. *Calif. Management Rev.* 43(4):89–105.
- Wübben M, Von Wangenheim F (2008) Instant customer base analysis: Managerial heuristics often “get it right.” *J. Marketing* 72(3):82–93.